

— Machine Learning Notes —

Contents

1	Introduction	1
1.1	Math	1
1.1.1	Eigenvalues and Eigenvectors	1
1.1.2	Derivative and Hessian	1
1.1.3	Bonus	2
2	Supervised learning	3
2.1	Least squares regression	3
2.2	Gradient descent	3
2.2.1	Derivative examples	4
2.2.2	Convexity	4
2.2.3	Backtracking line search	5
2.2.4	Solve LSR	5
2.2.5	Subgradient method	5
2.3	Polynomial Regression	6
2.4	Underfitting / Overfitting	6
2.4.1	k-fold Cross Validation	6
2.4.2	Regularization	7
2.4.3	Bias-Variance Tradeoff	7
2.4.4	Regularizers	7
2.5	Feature Scaling	8
2.5.1	MLE and MAP	8

1 Introduction

ML: Tries to automate the process of **inductive inference**.

1. Deduction: Learning from rules
2. Induction: Learning from examples

1.1 Math

TODO: norms

TODO: determinant, trace, inverse

TODO: semidefinite, definite, indefinite

TODO: linear eq

TODO: inverse proof

1.1.1 Eigenvalues and Eigenvectors

Example: $f(w) = 0.5w^T M w$

- Hessian: $\nabla^2 f(w) = M = \begin{pmatrix} 1 & 0 \\ 0 & 3 \end{pmatrix}$
- Eigenvalues 1, 3
- Function along eigenvectors like $1x^2$ and $3x^2$

1.1.2 Derivative and Hessian

$$L(w) : \mathbb{R} \rightarrow \mathbb{R}^m \Rightarrow \nabla L(w) = \begin{pmatrix} \frac{\partial}{\partial w_1} L(w) \\ \frac{\partial}{\partial w_2} L(w) \\ \vdots \\ \frac{\partial}{\partial w_n} L(w) \end{pmatrix} \Rightarrow \nabla L(w) = \begin{pmatrix} \frac{\partial L_1}{\partial w_1} & \frac{\partial L_2}{\partial w_1} & \cdots & \frac{\partial L_m}{\partial w_1} \\ \frac{\partial L_1}{\partial w_2} & \frac{\partial L_2}{\partial w_2} & \cdots & \frac{\partial L_m}{\partial w_2} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial L_1}{\partial w_d} & \frac{\partial L_2}{\partial w_d} & \cdots & \frac{\partial L_m}{\partial w_d} \end{pmatrix}$$

$$\nabla^2 f(x) = \begin{pmatrix} \frac{\partial^2 f}{\partial x_1 \partial x_1}(x) & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_d}(x) \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_d \partial x_1}(x) & \cdots & \frac{\partial^2 f}{\partial x_d \partial x_d}(x) \end{pmatrix}, \quad x = \begin{pmatrix} x_1 \\ \vdots \\ x_d \end{pmatrix} \in \mathbb{R}^d$$

1.1.3 Bonus

Convex hull (V) for set of vectors V is smallest convex set containing V .

$$(V) = \left\{ \sum_{i=1}^m \lambda_i \cdot v_i \mid \lambda_i \geq 0, \sum_{i=1}^m \lambda_i = 1 \right\}$$

2 Supervised learning

- input X , output Y
- training data: $(x^{(i)}, y^{(i)})_{i=1..n} \subset X \times Y$
- Goal: learn $f : X \rightarrow Y$ for model class F on examples

2.1 Least squares regression

\tilde{X}, \tilde{w} are extended with bias:

$$\min_{\tilde{w}} \frac{1}{2} \|\tilde{X} \tilde{w} - y\|^2 \Rightarrow \min_w \frac{1}{2} \|Xw - y\|^2$$

Solve with gradient and set to zero:

$$\begin{aligned} L &= \frac{1}{2} \sum_{i=1}^n ((X_i^T w_i) - y_i)^2 \\ &= \frac{1}{2} \left(\sum_{i=1}^n (X_i^T w_i)^2 - 2(X_i^T w_i)y_i + y_i^2 \right) \end{aligned}$$

$$\begin{aligned} \nabla L &= \frac{\partial}{\partial w} \left(\frac{1}{2} \left(\sum_{i=1}^n (X_i^T w_i)^2 - 2(X_i^T w_i)y_i + y_i^2 \right) \right) \\ &= \frac{1}{2} \left(\sum_{i=1}^n 2(X_i^T X_i w_i) - 2(X_i^T)y_i \right) \\ &= \sum_{i=1}^n X_i^T X_i w_i - X_i^T y_i \\ &= X^T X w - X^T y \\ &= X^T (Xw - y) \end{aligned}$$

$$\nabla L = X^T (Xw - y) = 0 \Rightarrow (X^T X)w = X^T y \Rightarrow w = (X^T X)^{-1} X^T y$$

2.2 Gradient descent

Alternative to least squares regression. Algorithm:

1. Compute gradient $\nabla L(w) = X^T (Xw - y)$
2. Negative gradient shows to steepest descent
3. $w^{(t+1)} = w^{(t)} - \gamma^{(t)} \cdot \nabla L(w^{(t)})$

2.2.1 Derivative examples

- $L(w) = w_1^2 + w_2^2$
 $\Rightarrow \nabla L(w) = \begin{pmatrix} 2w_1 \\ 2w_2 \end{pmatrix}$
- $L(w) = \|w\|_2^2 = w^T w$
 $\Rightarrow \nabla L(w) = 2w$
- $L(w) = w^T A w$
 $\Rightarrow \nabla L(w) = A w + A^T w$
- $L(w) = \|Xw - y\|^2 = w^T X^T X w - y^T X w - w^T X^T y + y^T y$
 $\Rightarrow \nabla L(w) = 2X^T (Xw - y)$

2.2.2 Convexity

Set C convex if line between any two points of C in C . $\forall x, y \in C$ and $\lambda \in \mathbb{R}$ with $0 \leq \lambda \leq 1$:

$$\lambda x + (1 - \lambda)y \in C$$

Function $f: \mathbb{R}^d \rightarrow \mathbb{R}$ convex if (f) is a convex set and $\forall x, y \in (f)$, $\lambda \in \mathbb{R}$ with $0 \leq \lambda \leq 1$:

$$f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y)$$

Gradient descent returns global optimum for convex functions.

Optimization problem: $\min f(x), x \in X \subseteq \mathbb{R}^d$ has local minimizer $x^* \in X$ if $\exists \varepsilon > 0$ with:

$$\forall y \in X \text{ with } \|x^* - y\| \leq \varepsilon : f(x^*) \leq f(y)$$

Global minimizer if $f(x^*)$ is lowest of all optimizers.

Symmetric matrix A is positive semidefinite ($A \succcurlyeq 0$) if :

$$x^T A x \geq 0, \forall x$$

Positive definite ($A \succ 0$) if $\forall x \neq 0$

Symmetric matrix A is positive semidefinite iff all eigenvalues are ≥ 0 and positive definite iff all > 0 .

If function is one-dimensional: Convex if $f''(x) \geq 0$. If multidimensional: Convex if 2nd derivative is psd.

2.2.3 Backtracking line search

Algorithm:

1. Input: $x, \Delta x, \alpha \in (0, 0.5), \beta \in (0, 1)$
2. $t = 1$
3. while $f(x + t \Delta x) > f(x) + \alpha t \nabla f(x)^T \Delta x$:
4. $t = \beta t$

2.2.4 Solve LSR

1. $L(w) = \frac{1}{2} \|Xw - y\|_2^2$
2. $\nabla L(w) = X^T(Xw - y)$
3. $\nabla L(w) = X^T X$ is symmetric and psd

2.2.5 Subgradient method

If function not differentiable, e.g. $\|w\|_1$

- gradient is subgradient (convex hull of gradients)
- choose constant step length g
- $w^{(t+1)} = w^{(t)} - \gamma^{(t)} \cdot g$ with $\gamma^{(t)} = \frac{1}{\sqrt{t}}$
- find $g \in \mathbb{R}^d$ at $x \in (f)$ with:

$$f(y) \geq f(x) + g^T(y - x), \forall y \in (f)$$

2.3 Polynomial Regression

- $X \in \mathbb{R}^S, Y \in \mathbb{R}^S$
- $f(x) = w_d x^d + w_{d-1} x^{d-1} + \dots + w_1 x^1 + w_0$
- find best $w = (w_d, \dots, w_0) \in \mathbb{R}^{d+1}$
- loss function is squared loss: $l(y, \hat{y}) = \frac{1}{2}(y - \hat{y})^2$

With $\hat{y} = f(x^{(i)}) = \sum_{j=0}^d w_j (x^{(i)})^j = (\tilde{x}^{(i)})^T w$ rewrite as:

$$\begin{aligned} w^* &= \min_w \sum_{i=1}^n \frac{1}{2} (y^{(i)} - \hat{y}^{(i)})^2 \\ &= \min_w \sum_{i=1}^n \frac{1}{2} (y^{(i)} - (\tilde{x}^{(i)})^T w)^2 \end{aligned} \quad (\text{LSR})$$

Solve $\|Xw - y\|^2$ with Basis functions:

$$X = \begin{pmatrix} f_1(x^{(1)}) & f_2(x^{(1)}) & \dots & f_m(x^{(1)}) \\ f_1(x^{(2)}) & f_2(x^{(2)}) & \dots & f_m(x^{(2)}) \\ \vdots & \vdots & \ddots & \vdots \\ f_1(x^{(n)}) & f_2(x^{(n)}) & \dots & f_m(x^{(n)}) \end{pmatrix} \quad y = \begin{pmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(n)} \end{pmatrix}$$

2.4 Underfitting / Overfitting

Underfitting: Model too simple, degree low

Overfitting: Model too complex, degree high

Too high model complexity \rightarrow Higher training error

Lower polynomial degree or basis functions \rightarrow Lower model complexity

2.4.1 k-fold Cross Validation

Mitigate Overfitting: Split training data into k (usually 10) and pick one for **validation data**.

Train model on one training block, run on validation data and compute error. Repeat for all blocks and average.

2.4.2 Regularization

Constrain magnitude ($\|w\|_2, \|w\|_1$, etc.)

Lagrangian to remove constraint

$$\min_w \begin{matrix} L(w) \\ \text{st} \quad \|w\|_2^2 \leq t \end{matrix} \quad \rightarrow \quad \min_w L(w) + \frac{\lambda}{2} \|w\|_2^2$$

if $L(w) = \frac{1}{n} \sum_{i=1}^n l(y^{(i)}, \hat{y}^{(i)})$:

1. Empirical risk minimization (ERM): $\min_w L(w)$
2. Regularized risk minimization (RRM): $\min_w L(w) + \|w\|$

2.4.3 Bias-Variance Tradeoff

Prediction error is sum of variance and bias

- Variance spreads predictions around true value
- Bias puts predictions away from true value

With complexer model:

1. Test data has min somewhere
2. Bias gets lower
3. Variance gets higher

2.4.4 Regularizers

Ridge Regression: LSR with $\|w\|_2$ -regularizer:

$$\min_w \frac{1}{2n} \|Xw - y\|_2^2 + \frac{\lambda}{2} \|w\|_2^2$$

Least absolute shrinkage and selection operator (LASSO): $\|w\|_1$ -regularizer:

$$\min_w \frac{1}{2n} \|Xw - y\|_2^2 + \lambda \|w\|_1$$

Solved with subgrad method, performs feature selection.

Elastic Net: Combination of both

$$\min_w \frac{1}{2n} \|Xw - y\|_2^2 + \lambda \left(\alpha \|w\|_1 + \frac{1-\alpha}{2} \|w\|_2^2 \right)$$

Often used for gene expression data.

Robust Regression with $\|w\|_1$ -regularizer:

$$\min_w \frac{1}{n} \|Xw - y\|_1$$

Solved with subgrad method. Often used with Huber Loss for faster, simpler optimization.

2.5 Feature Scaling

- Features should be $[0, 1]$ or $[-1, 1]$
- Regularizer not invariant to scaling
- also on test data!

Normalize data: Center and scale each feature of data matrix $X_{i,j} = (x_j^{(i)})$

$$X_{:,j}^{\text{centered}} = X_{:,j} - \bar{x}_j = X_{:,j} - \frac{1}{n} \sum_{i=1}^n x_j^{(i)}$$

$$X_{:,j}^{\text{scaled}} = \frac{X_{:,j}^{\text{centered}}}{\|X_{:,j}^{\text{centered}}\|_2}$$

2.5.1 MLE and MAP

Example: For Coin-throw with $p(\text{head}) = \theta$: 3 heads, 7 tails. What is most likely θ ?

$$p(y^{(1)}, y^{(2)}, \dots, y^{(n)} | \theta) = \prod_i p(y^{(i)} | \theta) = \theta^3 (1 - \theta)^7$$

Maximum Likelihood Estimator (MLE): Find θ for max probability:

$$\max_{\theta} \theta^3 (1 - \theta)^7$$

Maximum A Posteriori (MAP): Find θ for max probability with prior:

$$\max_{\theta} \theta^3 (1 - \theta)^7 \cdot p(\theta | \text{observation})$$

with $p(\theta | \text{observation}) = \frac{p(\text{observation}|\theta) \cdot p(\theta)}{p(\text{observation})}$

Empirical risk min.	Maximum likelihood
Minimize	Maximize
Sum	Product
Risk / Loss function	Noise Distribution
l_1 -loss	Gaussian Distribution
l_2 -loss	Laplacian Distribution